

Verifiable Layered Architecture Visualizations

Invariants for Correct Level Computation

[Draft Version 3 – for co-author review]

Johannes Weigend

Weigend AM GmbH & Co.KG, Germany/Brazil
johannes.weigend@weigend.de

Michael Philippsen

Friedrich-Alexander-Universität Erlangen-Nürnberg, Germany
michael.philippsen@fau.de

Keywords: software architecture visualization, layout invariants, layered software city, strongly connected components, dependency analysis, software engineering tools

Abstract: Layered architecture visualizations communicate dependency structure through vertical position: elements that depend on others are placed higher. If the level-assignment pipeline is wrong, the picture can still look plausible while misleading the viewer: cyclic peer packages may be placed on different levels, a regular dependency may point against the intended layer direction, or a top-level wiring class may visually sink below the sub-package it orchestrates. Layout correctness should be the tool’s responsibility, not the reader’s—this paper documents why correct level computation is harder than it looks.

Our 2021 IVAPP paper, *A Layered Software City for Dependency Visualization* (Dashuber et al., 2021), introduced a city layout in which the X/Y position of class buildings follows dependency-derived levels, while building height along the Z axis encodes orthogonal metrics such as complexity or lines of code. S202 reimplements and validates the dependency-level pipeline of that layout as an open-source Java tool under Apache 2.0, together with a Layout Invariant Checker that verifies the layout after every analysis run.

Correct level computation requires two separations. First, the architectural analysis must distinguish class-level dependency depth from package-level dependency depth: both are computed as the longest path in the respective SCC-collapsed dependency graph (unweighted for classes, weighted by method-call counts for packages), and neither is derived from the other. Second, and independently of both, the visual position of an element inside its parent container is governed by a per-parent local layer index, computed on the sibling-only dependency graph closed within that parent’s subtree. The local layer index is what places *ArchitectureView* above the *rendering* sub-package it calls into, without disturbing the global dependency depths. Highly cyclic code bases are handled by a weight-based rank heuristic that classifies back edges explicitly rather than hiding them.

The Layout Invariant Checker verifies five postconditions. Two are implementation sanity checks on the global architecture levels: non-heuristic class dependencies must respect class-level direction (R1-algo), and the dominant direction in the weighted package graph must be reflected in package levels (R3). Two are layout-level statements that the rendered picture is honest: every class dependency must point downward in the visual rank derived from the local layer index along the parent chain (R1-visual), and packages forming a cyclic group in the filtered package dependency graph must share an architecture level (R2). Finally, R5 keeps cached edge classifications in sync with the current level and SCC state. The checker distinguishes implementation bugs from tolerated heuristic violations and emits a copy-paste-ready reproducer block for every finding.

We will demonstrate the checker on real code bases, tracing the path from a marked layout element to a failing reproducer test. In the 92-module *software-ekg-7* system, an R2 finding exposed a missing package-SCC equalization step and directly triggered a pipeline fix. On Minecraft Forge 1.19.2, separating analysis depth from local layout drops the visual upward-edge count from roughly 6 000 to 3 600 by eliminating the artefacts of a previous parent-filtered weight pipeline; the remaining edges are real architectural debt. No previous layered architecture visualization known to us has provided this degree of machine-checkable self-validation.

REFERENCES

- Dashuber, V., Philippsen, M., and Weigend, J. (2021). A layered software city for dependency visualization. In *Proceedings of the 12th International Conference on Information Visualization Theory and Applications (IVAPP)*, VISIGRAPP 2021, pages 15–26. SCITEPRESS. Best Paper Award.
- Eades, P., Lin, X., and Smyth, W. F. (1993). A fast and effective heuristic for the feedback arc set problem. *Information Processing Letters*, 47(6):319–323.
- Karp, R. M. (1972). Reducibility among combinatorial problems. In Miller, R. E. and Thatcher, J. W., editors, *Complexity of Computer Computations*, pages 85–103. Plenum Press.

Appendix A – Why Naive Single-Pass Does Not Work: Failure Modes

The Core Problem: Two Coordinate Worlds, Four Constraints

A correct layered layout has to satisfy two coordinate concerns at once that are easy to conflate but answer different questions:

- *Analysis (global)* – where does this element sit in the dependency stack of the analysed system? This is the longest-path depth in the dependency graph and is best understood per element type: classes are levelled on the class graph, packages on the weighted inter-package graph. Both are global numbers.
- *Layout (local)* – where should the box for this element sit inside its parent container? This is decided by the dependency relationships among the parent’s direct children only; refs that leave the parent are irrelevant to local position because the parent itself is positioned by the same logic one level up.

The two worlds are independent in principle and need to be kept independent in the implementation. Mixing them is the root cause of every failure mode below. Four invariants follow from the two-world separation:

1. **Class-level analysis constraint (R1-algo):** If $A \rightarrow B$ denotes a class dependency, then $\text{archLvl}(A) > \text{archLvl}(B)$ after heuristic back-edges are removed. This is a sanity check on the longest-path assignment in the class-dependency SCC-DAG.
2. **SCC constraint (R2):** Classes in a cyclic dependency group receive the same class architecture level; packages that form a cyclic group in the filtered package dependency graph share the same package architecture level.
3. **Package-order constraint (R3):** In the weighted inter-package dependency graph, when the call traffic from package P_A to P_B dominates the reverse, P_A ends up at a strictly higher package architecture level than P_B .
4. **Visual-rank constraint (R1-visual):** For every class dependency $A \rightarrow B$ that is not a heuristic back-edge, the rendered position of A in its parent box chain sits above B . The comparison is lexicographic on the tuple of local layer indices along the parent path, ending with the class’s own local index. This is the statement the reader actually checks against the drawing.

R1-algo and R3 guard the algorithm; R1-visual and R2 guard the picture. The failure modes that follow show why each obvious shortcut collapses one of these into the other and loses information.

Failure Mode 1: Cascading Level Updates

An algorithm iterates over classes and assigns levels. It first sees class A in package P_1 , computes $\text{level}(A) = 3$, and therefore sets $\text{level}(P_1) = 3$. Later it reaches class B in package P_2 and discovers that A and B are in the same SCC and must share a level, now $\text{level}(A) = \text{level}(B) = 5$. This forces P_1 to be recomputed. If P_1 depends on P_3 , P_3 may need to move as well. A cascade follows.

In a naive single-pass implementation, however, P_1 and P_3 may already have been finalized. SCC membership is only known after the relevant graph has been traversed; if level assignment and SCC discovery are interleaved, the iteration order decides whether the correction is still applied in time.

Failure Mode 2: Interleaved Cycle Detection Produces Classpath-Dependent Results

The most subtle practical failure mode occurs when cycle detection and level computation are not separated. The algorithm traverses the class graph, assigns levels, and handles discovered cycles on the fly, depending on which class happens to be processed first.

Classes come from JARs. The order in which the JVM sees classes follows the order in which they were packaged into the JAR. That order can depend on the build tool, module order in a multi-module build, filesystem ordering on the build server, and other implementation details. None of these orders is semantically meaningful; none encodes architectural dependency direction.

The result is unreliable: the same source code can produce different level layouts under different build configurations. The errors are subtle. The layout may still look plausible, arrows may mostly point in the expected

direction, and no obvious visual break appears. Yet the precise hierarchy information is corrupted and may remain unnoticed for a long time.

The solution is to separate the phases strictly: first compute the SCC partition on the complete graph, then build the SCC-DAG, then assign levels. Tarjan's algorithm computes the SCC partition in $O(V + E)$; with deterministic input ordering, the implementation becomes reproducible and independent of classpath traversal artifacts.

Formal Counterexample: Why a Naive Recursive Algorithm Fails

The intuitive level-computation algorithm is a recursive DFS with memoization:

```
level(A) :
  if A already computed: return level(A)
  if A has no dependencies: return 0
  return max(level(dep) for dep in A.dependencies) + 1
```

For acyclic graphs this is correct and equivalent to longest-path computation on a DAG. For cyclic graphs it fails structurally.

Counterexample with two SCCs:

```
SCC_1 = {A, B}, with A -> B -> A
SCC_2 = {C, D}, with C -> D -> C
Cross dependency: A -> C
```

The correct levels are $\text{level}(C) = \text{level}(D) = 0$ and $\text{level}(A) = \text{level}(B) = 1$.

Recursive algorithm, starting at A:

```
level(A)  [A "in progress", sentinel = 0]
-> level(B)
    -> level(A): cycle detected, return 0
    B.level = 0 + 1 = 1
-> level(C)  [C "in progress", sentinel = 0]
    -> level(D)
        -> level(C): cycle detected, return 0
        D.level = 0 + 1 = 1
    C.level = 1 + 1 = 2    <-- wrong; correct would be 0
A.level = max(1, 2) + 1 = 3    <-- wrong; correct would be 1
```

After SCC equalization, $\text{SCC}_1 = \max(3, 1) = 3$ and $\text{SCC}_2 = \max(2, 1) = 2$. Both results are wrong.

The cause is that the recursive algorithm enters SCC_2 from the context of SCC_1 . The sentinel-based local cycle handling inflates the level of SCC_2 , and that inflated level then feeds into the computation of SCC_1 . The correct level of SCC_2 , namely zero, is only known after the graph has been collapsed into SCCs and evaluated as an SCC-DAG.

Thus, recursive level algorithms that mix traversal, cycle handling, and level assignment can produce traversal-order-dependent results on graphs with interacting SCCs. In JVM-based tools, that traversal order can be influenced by classpath and packaging order. The correct solution collapses SCCs before level assignment: Tarjan, then SCC-DAG, then longest path.

Failure Mode 3: Deriving Package Levels from Class Levels

An obvious implementation assigns each package a level equal to the maximum level of its contained classes, then lifts parent packages transitively. This is wrong for two reasons.

First, the rule conflates containment with dependency. A package that *contains* a class at level 7 is not necessarily architecturally positioned at level 7 itself; it may have no cross-package dependencies at all. Two disconnected package trees processed from the same JAR would be ordered relative to each other by accident of their class contents, not by their actual dependency relationships.

Second, the approach creates a circular dependency between the two level systems: class levels depend on the class graph, but package levels derived from class levels then constrain the visual position of packages relative to classes, making the combined coordinate system ill-defined.

S202 avoids this by computing class levels and package levels on independent graphs with no mutual constraint. Each system starts at zero and grows according to its own dependency structure.

Failure Mode 4: The Big Lake – SCCs Without Heuristics Collapse the Layout

Without a mechanism for breaking large SCCs, all members of a cyclic group must be placed on the same level. For small cycles of two to five classes this is acceptable. For Minecraft Forge 1.19.2 (5 123 classes), however, we observed a single SCC containing 2 857 classes. Without a breaker, more than 55 percent of the project would be assigned the same level. All remaining classes would be distributed above, but the internal dependency hierarchy within this large cycle would be invisible.

The heuristic SCC breaker solves this practical problem. It identifies back edges—dependencies that run against the natural architectural flow—and removes them for level computation. The removed edges are visualized as violations instead of being ignored. The layout preserves visibility of cyclic dependencies while still producing a usable hierarchy.

Relation to Feedback Arc Set Heuristics

Failure Mode 4 is an instance of the Minimum Feedback Arc Set problem: remove a minimal set of directed edges so the graph becomes acyclic. The problem is NP-hard (Karp, 1972), so any practical polynomial-time implementation must use a heuristic or approximation.

A well-known related heuristic is due to Eades, Lin, and Smyth (Eades et al., 1993). It orders nodes according to the difference between outgoing and incoming degree, then classifies edges that point backward in the resulting sequence as feedback edges. The intuition is similar to S202: high out-degree suggests a higher-level element; high in-degree suggests a foundational element.

S202 uses a deliberately simpler variant: a direct normalized score per node,

$$r(P) = \frac{\text{outWeight}(P) - \text{inWeight}(P)}{\max(1, \text{outWeight}(P) + \text{inWeight}(P))},$$

where the weights are summed within the SCC under consideration. At the class level the weights are the in- and out-degrees in the class dependency graph; at the package level they are aggregated method-call counts. An edge $P_A \rightarrow P_B$ inside an SCC is classified as a back-edge candidate when $r(P_A) < r(P_B) - 0.1$. Equal-rank pairs represent symmetric bidirectional dependencies with no dominant architectural direction; they remain in the same SCC and are assigned the same level.

The important distinction is transparency. Many tools break cycles implicitly, but S202 classifies heuristic cut edges explicitly. They are marked as `VIOLATION`, rendered in red, and known to the invariant checker. This makes algorithmic defects separable from tolerated heuristic artifacts.

The Solution: Three Coordinate Concepts on Independent Graphs

Failure modes 1 to 3 share a root cause: conflating coordinate concepts that belong to different semantic domains. S202 resolves this by computing three independent coordinate values, two of them for analysis and one for layout.

- **Class architecture level** is the longest path in the SCC-collapsed class dependency DAG. It represents global architectural depth: how many layers of dependencies separate a class from the foundation of the analysed system.
- **Package architecture level** is the longest path in the SCC-collapsed weighted inter-package dependency graph. Edge weights are method-call counts aggregated from class-level references. The result represents the module's position in the system's overall dependency order, independently of the class-level coordinate.
- **Local layer index** is the longest path in the SCC-collapsed sibling-only dependency graph of a single parent container. The graph contains only edges whose source and target both lie in some direct child of that parent; refs leaving the parent's subtree are excluded because they are the business of the enclosing container. Every

element (class and package) gets a local layer index that is only meaningful within its own parent box. The renderer sorts each parent’s children by this index, so the visible Y position of a box is the local-layer-index world, not the global one.

All three concepts use the same algorithmic shape—Tarjan SCC detection, weight-based back-edge identification, DAG condensation, longest-path propagation—but on three different graphs with three different scopes. Because the systems are independent, packages with no cross-package dependencies correctly land at architecture level 0 regardless of the class levels inside them, and a top-level class like *ArchitectureView* that orchestrates a sub-package such as *ui.rendering* sits visually above the sub-package box without distorting either of their architecture levels.

The separation is an engineering choice driven by semantic clarity. A unified single-coordinate approach is theoretically possible, but it conflates three semantically distinct concerns—global architectural depth, module dependency order, and per-container visual position—into one number, making the pipeline hard to reason about, test, and maintain. Every obvious simplification of such a unified approach turns into one of the failure modes above.

The three concepts have the following algorithmic structure. Phases 1 and 2 compute the global architecture levels; Phase 3 derives the local layer index per parent container from the architecture-level output.

Algorithm 1: Phase 1 – Global Class Levels (*LevelCalculator*, Step 3)

```

class dependency graph
  → Tarjan (SCCs)
  → SCCBreaker: identify heuristic back-edges in large SCCs ( $|S| \geq 3$ ; 2-class SCCs are equalized instead)
     $\text{rank}(P) = (\text{outDeg} - \text{inDeg}) / \max(1, \text{outDeg} + \text{inDeg})$ 
    edge  $A \rightarrow B$  is a back-edge candidate if  $\text{rank}(A) < \text{rank}(B) - 0.1$ 
  → Remove back-edges → filtered class DAG
  → Condense into SCC-DAG → longest-path pass
  → Assign SCC level to all member classes (equalization implicit)

```

Algorithm 2: Phase 2 – Package Architecture Levels (*LevelCalculator*, Step 4)

Build weighted inter-package dependency graph $G_P = (V_P, E_P, w)$:

For each class C in package P_{leaf} with method calls to class D in P_{target} ($P_{\text{leaf}} \neq P_{\text{target}}$):

$$w(P_{\text{ancestor}}, P_{\text{target}}) += \text{callCount}(C \rightarrow D)$$

for every ancestor P_{ancestor} of P_{leaf} with $P_{\text{ancestor}} \neq P_{\text{target}}$

Fallback weight = 1 for structural references with no method-call data

No directional filtering: both parent→child and child→ancestor edges enter G_P . The SCC breaker decides direction by weight.

Iterative SCC-breaking (repeat until stable):

Tarjan on current graph → find SCCs with $|S| \geq 2$ (2-package SCCs broken if rank differs; equalized otherwise)

For each member $P \in S$:

$$r(P) = (\sum_{Q \in S} w(P, Q) - \sum_{Q \in S} w(Q, P)) / \max(1, \sum_{Q \in S} w(P, Q) + \sum_{Q \in S} w(Q, P))$$

Edge $P_A \rightarrow P_B$ inside S is a back-edge if $r(P_A) < r(P_B) - 0.1$; remove it

Equal-rank pairs remain in the SCC → assigned the same level

Level assignment:

- Condense remaining graph into SCC-DAG
- Longest-path propagation on SCC-DAG
- Assign SCC level to all member packages

Phase 1 produces classpath-independent class architecture levels because SCCs are collapsed before levels are assigned. Phase 2 produces package architecture levels honestly: every class reference contributes to the weighted graph regardless of the containment relation between source and target package, and a per-SCC rank score decides which direction of a cyclic group dominates. The two architecture coordinate systems do not constrain each other – a package with no outgoing cross-package calls lands at level 0 regardless of the class

levels inside it, and conversely the class levels of a package’s contents do not lift the package itself.

Algorithm 3: Phase 3 – Local Layer Index per Parent Container (`LocalLayerCalculator`, Step 6)

for each parent package P with at least two direct children:
 Let $S(P)$ be the direct children (classes and sub-packages).
 Build sibling-only weighted graph G_P^{loc} on $S(P)$:
 edge $X \rightarrow Y$ with weight $\sum \text{callCount}(c \rightarrow d)$
 over all class refs $c \rightarrow d$ where c lies in the subtree of sibling X
 and d lies in the subtree of sibling Y ;
 refs leaving the parent’s subtree are excluded.
 Tarjan on $G_P^{\text{loc}} \rightarrow \text{SCCs}$.
 Within each SCC of size ≥ 2 : weight-based rank break (same formula as Phase 2).
 Condense to SCC-DAG, longest-path pass, assign **localLayerIndex** per sibling.

Phase 3 is a pure layout phase. It never modifies architecture levels. Each parent container is a self-contained world: a class’s local layer index inside ui carries no relationship to a class’s local layer index inside $ui.rendering$. The renderer sorts each parent’s children by their local layer index, descending. Cross-container ordering is the business of the enclosing parent’s local layer index, applied recursively. Local cycles—two sibling sub-packages with mutual class references—are resolved silently by the rank break; they are not surfaced as a separate Tangle category because the user-visible tangle list remains a global architecture statement.

	Phase 1 – Class	Phase 2 – Package	Phase 3 – Local Layer
Graph	Class dependency	Weighted inter-package	Sibling-only weighted (per parent)
Edge weight	Unweighted	Method-call counts	Method-call counts
Scope	Global	Global	Per parent container
Meaning	Class arch. depth	Module dep. position	Visual Y in parent box
Zero point	Class with no deps	Package with no outgoing calls	Bottom sibling in box
Cycles	SCCBreaker (degree)	Weight-rank, all edges	Weight-rank within siblings
Used by	R1-algo, Tangles, R2 (class)	R3, R2 (package)	R1-visual, renderer sort

Table 1: The three coordinate concepts computed by S202.

Layout Invariants for Level Correctness

The three level phases define five machine-checkable postconditions verified after every analysis run. R1-algo, R2, R3, R5 are checked by the `LayoutInvariantChecker` as algorithm sanity guards; R1-visual is checked by the architecture-builder’s violation detection on the visual-rank tuple.

Rule	Postcondition
R1-algo	Non-heuristic class dependencies respect class architecture levels: $A \rightarrow B \Rightarrow \text{archLvl}(A) > \text{archLvl}(B)$. Tests the longest-path assignment of Phase 1.
R1-visual	For every non-heuristic class dependency $A \rightarrow B$, the rendered position of A sits above B . Position is compared lexicographically on $[L(\text{ancestor}_1), \dots, L(\text{ancestor}_n), L(A)]$ versus $[L(\text{ancestor}'_1), \dots, L(\text{ancestor}'_m), L(B)]$ where L denotes the local layer index. Covers same-parent, same-grandparent, and cross-tree comparisons uniformly.
R2	Packages forming a cyclic group in the filtered package dependency graph share the same package architecture level.
R3	For every weighted package edge (P_A, P_B) with $w(P_A, P_B) > w(P_B, P_A)$ that is not a back-edge cut by the SCC breaker: $\text{archLvl}(P_A) > \text{archLvl}(P_B)$.
R5	Cached edge classifications (<code>NORMAL</code> , <code>VIOLATION</code> , <code>BACK_EDGE</code>) remain consistent with the current level assignment and SCC membership.

Table 2: Five layout invariants verified after every analysis run. R1-algo, R3 and R2 guard the algorithm; R1-visual guards the rendered picture; R5 keeps the cached annotations honest.

Throughout the iterative development of the level pipeline, R2 surfaced implementation bugs after algorithm changes that the unit-test suite had not detected. Figure 1 shows a representative R2 report. In *software-ekg-7*, for instance, an R2 finding exposed a missing package-SCC equalization step and directly triggered a pipeline fix.




Figure 1: R2 checker report: packages in a filtered package-SCC are placed on different levels. Findings are reported as algorithm-bug candidates with a copy-paste-ready reproducer block, which makes the visual failure traceable to concrete level-pipeline constraints.

R2 also guided refinement of the rule itself. Against Minecraft Forge (2 895 classes, 220 packages, 188 of 207 packages in a single package-level SCC), two R2 reports turned out to be *false positives*: R2 filtered class-level heuristic back-edges but not the package-level back-edges that the package SCC-breaker had already removed. R3 already applied the correct filter; once R2 was aligned, the Minecraft findings disappeared and all 155 invariant tests pass. The invariant checker connects the visual symptom to a reproducible algorithm defect.

The Remaining Rules: R1-algo, R1-visual, R3, and R5

R1-algo ($A \rightarrow B \Rightarrow \text{archLvl}(A) > \text{archLvl}(B)$ on the filtered class graph) is the most direct invariant on the analysis side: it is a corollary of longest-path computation on the SCC-DAG and holds trivially for any correctly implemented class-level assignment. Its value lies in providing an immediate sanity check—a failing R1-algo signals a fundamental breakdown in the class-level pipeline.

R1-visual (lexicographic comparison of local-layer-index tuples along the parent chain) is the invariant the reader actually checks against the drawing. The two checks are intentionally separated: R1-algo asks “did Phase 1 compute consistent class architecture levels?”, R1-visual asks “does the rendered Y position honestly reflect the class dependencies?” Both pass on a correct implementation; only the second one is allowed to fail when the sibling-only graph of some parent exposes a genuine upward edge, because that is real architectural debt rather than an algorithm bug. The single visual-rank comparison covers every pair: siblings inside the same container compare on their final tuple element; cross-tree pairs diverge earlier; classes versus sub-package boxes in the same parent compare on the shared parent context. No special-case logic is needed for parent→child versus sibling versus cross-package class references.

R3 verifies that the dominant direction in the weighted package dependency graph is reflected in the architecture level assignment. For any package edge (P_A, P_B) where $w(P_A, P_B) > w(P_B, P_A)$ —meaning P_A calls into P_B significantly more than vice versa—the invariant requires $\text{archLvl}(P_A) > \text{archLvl}(P_B)$. Package back-edges identified during SCC-breaking are excluded from this check, analogously to how R1-algo excludes class-level heuristic back-edges. A failing R3 indicates that the weight-based SCC-breaking cut the wrong package edge, or that the package-level DAG propagation did not converge correctly.

R5 validates that every edge’s stored classification (NORMAL, VIOLATION, BACK_EDGE) is still consistent with the current level assignment and SCC membership. Edge labels are computed during analysis and cached in the model; after any algorithmic change they can silently become stale. R5 is therefore a cross-cutting consistency guard: it ties the structural output of the level pipeline back to the semantic annotations visible to the user.

R4: A Retired Rule and What It Teaches

An early version of the invariant checker contained R4:

`Cyclic child packages must share a level.`

The rule sounds reasonable and is correct in simple cases. In practice it failed because the raw package graph contains relationships that should not be interpreted as peer-level architectural cycles:

- **Parent-child edges** create apparent bidirectional relationships because a parent contains its children and children belong to their parent. Without filtering, many parent/child pairs look cyclic.
- **SCCBreaker back edges** connect packages that are not necessarily architectural peers. A rule that includes these edges equalizes the wrong candidates.
- **Shared class SCCs** can make two packages appear package-dependent without representing a true structural peer relationship.

R4 was replaced by R2, which expresses the same intention on the correctly filtered graph. R2 removes class-level heuristic back-edges, package-level SCC-breaking back-edges, intra-subtree edges, and edges caused by shared class SCCs before running Tarjan on the remaining package graph.

Specifying invariants for a three-level constraint system turns out to be harder than it appears. The checker therefore provides a second safety net: it not only tests the implementation against the rules, but also exposes when the rules themselves are too broad or too coarse. False positives force the rule to become more precise. The later R2 finding in *software-ekg-7* shows that the refined rule was precise enough to distinguish real pipeline bugs from apparent ones.

Implementation: The Complete Pipeline

The conceptual structure expands into the following concrete implementation steps, which constitute the reference implementation (Appendix B).

Algorithm 1: Complete Level Pipeline (LevelCalculator)

- Step 1:** Create class objects (all levels $\leftarrow 0$)
- Step 2:** Create package objects (all levels $\leftarrow 0$)
- Step 3:** Compute class architecture levels (Phase 1)
- Tarjan on class dependency graph
 - SCCBreaker: identify heuristic back-edges in large SCCs ($|S| \geq 3$):
 $r(C) = (\text{outDeg} - \text{inDeg}) / \max(1, \text{outDeg} + \text{inDeg})$;
edge $A \rightarrow B$ is a back-edge candidate if $r(A) < r(B) - 0.1$
 - Remove identified back-edges → filtered class DAG
 - Condense into SCC super-nodes
 - Single longest-path pass on condensed DAG
 - Assign SCC level to all member classes
- Step 4:** Compute package architecture levels (Phase 2)
- 4a.** Build weighted inter-package dependency graph:
 $w(P_A, P_B) = \sum \text{method-call counts from subtree}(P_A) \text{ to } P_B$;
propagate weights up to all ancestor packages (no directional filtering);
both parent→child and child→ancestor edges enter the graph
- 4b.** Iterative package SCC-breaking (repeat until stable):
Tarjan on current graph → SCCs with $|S| \geq 2$;
 $r(P) = (\sum_{Q \in S} w(P, Q) - \sum_{Q \in S} w(Q, P)) / \max(1, \dots)$;
edge $P_A \rightarrow P_B$ is a back-edge if $r(P_A) < r(P_B) - 0.1$; remove it;
equal-rank pairs remain in the same SCC → assigned the same level
- 4c.** Level assignment on condensed package SCC-DAG:
longest-path propagation; assign SCC level to all member packages.
No child–parent lift: layout positioning is the local-layer-index world.
- Step 5:** Set reverse dependencies (dependents)
- Step 6:** Compute local layer index per parent (Phase 3)
- for each** parent package P with two or more direct children:
- build sibling-only weighted graph G_P^{loc} from class refs
whose source and target both live in some direct child of P ;
 - Tarjan → rank-based SCC break → longest-path;
 - assign `localLayerIndex` to each direct child of P .
-

After Step 3, a class’s architecture level is the level of the SCC node to which it belongs, measured as the longest path in the cycle-free quotient graph across all classes of the analysed project. Cycles do not influence level assignment internally because they have been collapsed before longest-path computation begins.

Step 4 treats class architecture levels and package architecture levels as independent coordinate systems. A package with no outgoing cross-package calls receives architecture level 0, regardless of the class levels inside it; two package trees with no shared dependencies are each levelled from zero in their own connected component of the weighted package graph. The earlier child–parent lift, which inflated parent package levels by the class levels their children called, is gone – it was a workaround for the missing layout coordinate.

Step 6 introduces the layout coordinate. Each parent container is treated as a self-contained universe: only class refs that stay within the parent’s subtree contribute to the local sibling graph. The result is that the box for *ArchitectureView* in the *ui* package sorts above the box for the *ui.rendering* sub-package, simply because *ArchitectureView* has heavier outgoing call traffic into rendering than rendering has back into *ui*. None of this disturbs Phase 1 or Phase 2.

Algorithm 2: Phase 4 – Horizontal Row Ordering (HorizontalRowLayoutOptimizer)

for each package P :

Step 1: Group direct children by their `localLayerIndex` value

Step 2: Collect subtree classes per row element

Step 3: Build weighted undirected proximity links from inter-sibling class dependencies

Step 4: Run bounded barycentric sweeps across rows


Step 5: Store the result in `horizontalLayoutOrder`

rendering: sort a copied child view by `localLayerIndex` (Phase 3), then `horizontalLayoutOrder` (Phase 4), then full name

Phase 4 does not change any level or layer assignment and is deliberately excluded from the layout invariants. Its only purpose is cosmetic: when an overlay of dependency arrows is displayed, same-row elements that share many inter-sibling calls are placed horizontally near each other so the arrows are shorter and less crossed. The implementation stores this as a separate `horizontalLayoutOrder` field and never mutates the semantic child list.

R2 Fix in *software-ekg-7*: Visual Evidence


The *software-ekg-7* case is the concrete failure that motivated the package-level pipeline. Before the fix, the city view still looked plausible: the package rectangles and dependency streets formed a readable layout, and no single visual element made the pipeline defect obvious. The metrics panel, however, exposed the inconsistency: 167 dependency violations were reported and the maximum architectural level had inflated to 160. After the package-level equalization was corrected, the same input produced zero violations and a maximum level of 18. The number of classes, packages, dependencies, cyclic classes, and tests stayed constant; the change is therefore attributable to level propagation, not to a different input model.



(a) Before the fix: the full city view looks plausible, but the metric panel reports 167 violations and MaxLevel 160.

(b) After the fix: the same analyzed system reports zero violations and MaxLevel 18.

Figure 2: Full-screen before/after screenshots of the *software-ekg-7* software city. The visible city changes only subtly; the level metrics reveal the correction.



(a) Before: 167 violations, MaxLevel 160.

(b) After: zero violations, MaxLevel 18.

Figure 3: Cropped metric panels from Figure 2. These panels are the clearest user-visible symptom of the propagation defect and its fix.

Summary

Correct level computation for layered architecture visualizations requires keeping two coordinate worlds separate: *analysis* (where does this element sit in the global dependency stack?) and *layout* (where should its box go inside its parent container?). Every straightforward single-pass approach over a unified graph collapses these two questions into one coordinate, and the failure modes documented above show why each obvious simplification breaks down in practice.

S202 resolves this by computing three coordinate values on three independent graphs. *Class architecture levels* are the longest path in the SCC-collapsed class dependency DAG; they represent global architectural depth. *Package architecture levels* are derived from a weighted inter-package dependency graph whose edges carry aggregated method-call counts; they represent module dependency position independently of class levels. *Local layer indices* are computed per parent container on the sibling-only graph closed within that parent’s subtree; they decide visual Y position inside each box without disturbing either architecture coordinate. All three use the same algorithmic shape: Tarjan SCC detection, weight-based back-edge identification, DAG condensation, longest-path propagation.

A horizontal row-ordering pass (Phase 4) improves the legibility of dependency overlays without modifying any coordinate or affecting any layout invariant. The Layout Invariant Checker verifies five machine-checkable postconditions (R1-algo, R1-visual, R2, R3, R5) after every analysis run and emits copy-paste-ready reproducer blocks for each finding.

Two cases demonstrate the value of the separation. In *software-ekg-7*, a plausible city layout reported 167 dependency violations and a maximum architectural level of 160 before a package-level equalization fix; after the fix, the same input produced zero violations and a maximum level of 18 with no change to the analysed source code. On Minecraft Forge 1.19.2, introducing the local layer index as a distinct layout coordinate and lifting the parent filter from the weighted package graph reduced the visual upward-edge count from roughly 6 000 to 3 600; the eliminated edges were artefacts of forcing parent packages below their sub-packages to keep them out of the algorithm’s tangle list, and the remaining edges expose real architectural debt that deserves a code-level fix.

Appendix B – LevelCalculator Reference Implementation

The following source excerpt is the complete Java implementation of the level pipeline discussed above.

```
1 package de.weigend.s202.domain.architecture;
2
3 import de.weigend.s202.domain.DomainModel;
4 import de.weigend.s202.domain.strategy.LevelCalculationStrategyContext;
5 import de.weigend.s202.graph.StronglyConnectedComponent;
6 import de.weigend.s202.graph.TarjanSCCFinder;
7 import de.weigend.s202.reader.DependencyModel;
8
9 import java.util.*;
10 import java.util.logging.Logger;
11
12 /**
13  * Calculates architectural levels for classes and packages based on dependencies.
14  *
15  * Pipeline:
16  *   Step 1 Create class objects (all levels = 0)
17  *   Step 2 Create package objects (all levels = 0)
18  *   Step 3 Compute class levels strategy (Tarjan → SCC-break → DAG → longest-path)
19  *   Step 4 Compute package levels weighted inter-package graph → SCC-break → DAG → longest-path
20  *   + child→parent lift for class-level alignment
21  *   Step 5 Set reverse dependencies
22  *   Step 6 Assign local layer index per-parent sibling graph → SCC-break → DAG → longest-path
23  *   (rendering position within each parent box, no global meaning)
24  *
25  * Package levels are computed independently from class levels using a weighted
26  * inter-package dependency graph. The weight of an edge P_A → P_B is the number
27  * of distinct classes in P_A that depend on at least one class in P_B (intra-subtree
28  * edges excluded). This separates the containment hierarchy from the dependency
29  * hierarchy and correctly handles disconnected package trees (they are levelled
30  * independently, both starting at 0).
31  *
32  * Package SCC-breaking uses a weight-based rank score:
33  *   rank(P) = (Σ outgoing weights within SCC – Σ incoming weights within SCC)
34  *             / max(1, sum of both)
35  * Edge P_A→P_B is a back-edge when rank(P_A) < rank(P_B) – 0.1.
36  * Equal-rank pairs (symmetric dependency) are left in the same SCC and assigned
37  * the same level -- they are genuine peers with no dominant dependency direction.
38  */
39 public class LevelCalculator {
40
41     private static final Logger LOG = Logger.getLogger(LevelCalculator.class.getName());
42     private static final double RANK_THRESHOLD = 0.1;
43
44     private final LevelCalculationStrategyContext strategyContext;
45
46     public LevelCalculator() {
47         this(LevelCalculationStrategyFactory.createDefault());
48     }
49
50     public LevelCalculator(LevelCalculationStrategyContext strategyContext) {
51         Objects.requireNonNull(strategyContext, "strategyContext cannot be null");
52         this.strategyContext = strategyContext;
53     }
54
55     public DomainModel calculate(DependencyModel rawModel) {
56         DomainModel model = new DomainModel();
57
58         // Step 1: Create class objects (all levels = 0)
59         for (String className : rawModel.getAllClassNames()) {
```

```

60     DependencyModel.ClassInfo rawClass = rawModel.getClass(className);
61     model.addClass(className, new DomainModel.CalculatedElementInfo(
62         className, rawClass.simpleName, "CLASS", 0,
63         new HashSet<>(rawClass.dependencies), rawClass.interfaceType));
64 }
65
66 // Step 2: Create package objects (all levels = 0)
67 for (String packageName : rawModel.getAllPackageNames()) {
68     DependencyModel.PackageInfo rawPkg = rawModel.getPackage(packageName);
69     model.addPackage(packageName, new DomainModel.CalculatedElementInfo(
70         packageName, rawPkg.simpleName, "PACKAGE", 0, new HashSet<>()));
71 }
72
73 // Step 3: Compute class levels
74 calculateClassLevels(model);
75
76 // Step 4: Compute package levels from weighted inter-package graph.
77 calculatePackageLevels(model, rawModel);
78
79 // Step 5: Set reverse dependencies
80 updateDependentRelationships(model);
81
82 // Step 6: Assign per-parent local layer index -- independent of the
83 // global architectureLevel, used by the renderer to position
84 // siblings within each parent's box.
85 new LocalLayerCalculator().assign(model, rawModel);
86
87 return model;
88 }
89
90 public LevelCalculationStrategyContext getStrategyContext() {
91     return strategyContext;
92 }
93
94 // -----
95 // Step 3 -- class levels
96 // -----
97
98 private void calculateClassLevels(DomainModel model) {
99     Map<String, Set<String>> classDeps = new HashMap<>();
100     for (DomainModel.CalculatedElementInfo c : model.getAllClasses().values()) {
101         classDeps.put(c.fullName, new HashSet<>(c.dependencies));
102     }
103     Map<String, Integer> levels =
104         strategyContext.getClassLevelStrategy().calculateClassLevels(classDeps);
105     for (Map.Entry<String, Integer> e : levels.entrySet()) {
106         DomainModel.CalculatedElementInfo info = model.getClass(e.getKey());
107         if (info != null) info.setArchitectureLevel(e.getValue());
108     }
109 }
110
111 // -----
112 // Step 4 -- package levels via weighted inter-package graph
113 // -----
114
115 private void calculatePackageLevels(DomainModel model, DependencyModel rawModel) {
116     if (model.getAllPackages().isEmpty()) return;
117     Set<String> allPkgNames = model.getAllPackages().keySet();
118
119     // Build weighted graph: weight[from][to] = # distinct classes in 'from'
120     // with at least one dependency on a class in 'to' (subtree edges excluded).
121     Map<String, Map<String, Integer>> weights = buildWeightedPackageGraph(model, rawModel);

```

```

122
123 // Unweighted adjacency for Tarjan (direction matters, zero-weight edges excluded)
124 Map<String, Set<String>> graph = new LinkedHashMap<>();
125 for (String pkg : allPkgNames) graph.put(pkg, new LinkedHashSet<>());
126 for (Map.Entry<String, Map<String, Integer>> entry : weights.entrySet()) {
127     graph.get(entry.getKey()).addAll(entry.getValue().keySet());
128 }
129
130 // Iteratively break SCCs using weight-based rank, then assign levels
131 // via SCC-collapsed DAG longest-path. Remaining SCCs (equal-rank peers)
132 // are collapsed to the same level without cutting any edge.
133 Set<String> backEdgeKeys = new LinkedHashSet<>(); // "from\Oto" -- tracked for R3
134 boolean changed = true;
135 while (changed) {
136     changed = false;
137     for (StronglyConnectedComponent scc : new TarjanSCCFinder(graph).findSCCs()) {
138         if (scc.getSize() < 2) continue;
139         Set<String> members = scc.getMembers();
140
141         // Compute weight-based rank for each member within this SCC
142         Map<String, Double> rank = new HashMap<>();
143         for (String m : members) {
144             int out = 0, in = 0;
145             for (String other : members) {
146                 if (other.equals(m)) continue;
147                 out += weights.getOrDefault(m, Map.of()).getOrDefault(other, 0);
148                 in += weights.getOrDefault(other, Map.of()).getOrDefault(m, 0);
149             }
150             rank.put(m, (out - in) / (double) Math.max(1, out + in));
151         }
152
153         // Identify back-edges: from→to where rank(from) < rank(to) - threshold
154         for (String from : new ArrayList<>(members)) {
155             for (String to : new ArrayList<>(graph.getOrDefault(from, Set.of()))) {
156                 if (!members.contains(to)) continue;
157                 if (rank.get(from) < rank.get(to) - RANK_THRESHOLD) {
158                     String key = from + "\0" + to;
159                     if (backEdgeKeys.add(key)) {
160                         graph.get(from).remove(to);
161                         changed = true;
162                     }
163                 }
164             }
165         }
166     }
167 }
168
169 // Assign package levels: SCC-collapsed DAG → longest-path
170 List<StronglyConnectedComponent> sccs = new TarjanSCCFinder(graph).findSCCs();
171 Map<String, StronglyConnectedComponent> pkgToScc = new HashMap<>();
172 for (StronglyConnectedComponent scc : sccs) {
173     for (String m : scc.getMembers()) pkgToScc.put(m, scc);
174 }
175
176 // SCC dependency graph (between SCC IDs)
177 Map<Integer, Set<Integer>> sccDeps = new HashMap<>();
178 for (StronglyConnectedComponent scc : sccs) {
179     sccDeps.put(scc.getId(), new HashSet<>());
180     for (String m : scc.getMembers()) {
181         for (String to : graph.getOrDefault(m, Set.of())) {
182             StronglyConnectedComponent toScc = pkgToScc.get(to);
183             if (toScc != null && toScc.getId() != scc.getId()) {

```

```

184         sccDeps.get(scc.getId()).add(toScc.getId());
185     }
186 }
187 }
188 }
189
190 // Longest-path levels on the SCC-collapsed DAG. The previous
191 // "childPkgUsedLevel" lift (which inflated a parent's level by
192 // the class levels its child packages used) is gone -- layout
193 // positioning lives on localLayerIndex now, so architectureLevel
194 // can be an honest dep-chain depth.
195 Map<Integer, Integer> sccLevels = new HashMap<>();
196 for (StronglyConnectedComponent scc : sccs) sccLevels.put(scc.getId(), 0);
197 boolean lvlChanged = true;
198 while (lvlChanged) {
199     lvlChanged = false;
200     for (StronglyConnectedComponent scc : sccs) {
201         int maxDep = -1;
202         for (int depId : sccDeps.getOrDefault(scc.getId(), Set.of())) {
203             maxDep = Math.max(maxDep, sccLevels.getOrDefault(depId, 0));
204         }
205         int newLevel = maxDep >= 0 ? maxDep + 1 : 0;
206         if (sccLevels.get(scc.getId()) != newLevel) {
207             sccLevels.put(scc.getId(), newLevel);
208             lvlChanged = true;
209         }
210     }
211 }
212
213 // Apply levels to all packages
214 Map<String, DomainModel.CalculatedElementInfo> pkgInfos = model.getAllPackages();
215 for (StronglyConnectedComponent scc : sccs) {
216     int level = sccLevels.get(scc.getId());
217     for (String member : scc.getMembers()) {
218         DomainModel.CalculatedElementInfo pkg = pkgInfos.get(member);
219         if (pkg != null) pkg.setArchitectureLevel(level);
220     }
221 }
222
223 // Store the weighted graph and identified back-edges in the model
224 model.setPackageEdgeWeights(weights);
225 model.setPackageBackEdges(backEdgeKeys);
226
227 // Populate package dependencies (unweighted) for reverse-dependency tracking
228 for (Map.Entry<String, Map<String, Integer>> entry : weights.entrySet()) {
229     DomainModel.CalculatedElementInfo pkg = pkgInfos.get(entry.getKey());
230     if (pkg != null) entry.getValue().keySet().forEach(pkg::addDependency);
231 }
232 }
233
234 /**
235  * Builds the weighted inter-package dependency graph.
236  * weight( $P_A \rightarrow P_B$ ) = total method-call count from classes in  $P_A$ 's subtree
237  * to classes in  $P_B$ . Method calls are a stronger signal than distinct-class
238  * counts: a class called hundreds of times clearly dominates one called once,
239  * making SCC-breaking direction unambiguous. Intra-subtree calls are excluded.
240  * Child-to-ancestor edges (e.g. sub-package calling into its parent package)
241  * are included; parent-to-child edges are excluded.
242  * Each call count is propagated to all ancestor packages up to the point where
243  * the ancestor would enter the same subtree as the target.
244  */
245 private Map<String, Map<String, Integer>> buildWeightedPackageGraph(

```



```

246     DomainModel model, DependencyModel rawModel) {
247     Map<String, Map<String, Integer>> weights = new HashMap<>();
248     for (String pkg : model.getAllPackages().keySet()) weights.put(pkg, new LinkedHashMap<>());
249
250     Set<String> allPkgNames = weights.keySet();
251
252     for (DomainModel.CalculatedElementInfo cls : model.getAllClasses().values()) {
253         String leafPkg = extractPackageName(cls.fullName);
254         if (leafPkg == null || !allPkgNames.contains(leafPkg)) continue;
255
256         // Count method calls per target package using raw model data
257         Map<String, Integer> callCountPerPkg = new LinkedHashMap<>();
258         DependencyModel.ClassInfo rawCls = rawModel.getClass(cls.fullName);
259         if (rawCls != null) {
260             for (DependencyModel.MethodInfo method : rawCls.methods.values()) {
261                 for (Map.Entry<String, Integer> call : method.methodCalls.entrySet()) {
262                     // Method call key format: "ownerClass.methodName"
263                     // Find the target class by checking known dependencies
264                     String calledKey = call.getKey();
265                     for (String dep : cls.dependencies) {
266                         if (calledKey.startsWith(dep + ".")) {
267                             String toPkg = extractPackageName(dep);
268                             if (toPkg != null && !leafPkg.equals(toPkg)
269                                 && allPkgNames.contains(toPkg)) {
270                                 callCountPerPkg.merge(toPkg, call.getValue(), Integer::sum);
271                             }
272                             break;
273                         }
274                     }
275                 }
276             }
277         }
278         // Fall back to 1 per target package for structural deps with no call data
279         for (String dep : cls.dependencies) {
280             String toPkg = extractPackageName(dep);
281             if (toPkg != null && !leafPkg.equals(toPkg) && allPkgNames.contains(toPkg)) {
282                 callCountPerPkg.putIfAbsent(toPkg, 1);
283             }
284         }
285
286         // Propagate each weight up to ancestor packages. Parent->child
287         // edges are no longer filtered here -- the rank-based SCC breaker
288         // decides direction based on actual call-count weight. Layout
289         // positioning is the LocalLayerCalculator's job (step 6), so the
290         // global package level can finally be an honest dep-chain count.
291         for (Map.Entry<String, Integer> entry : callCountPerPkg.entrySet()) {
292             String toPkg = entry.getKey();
293             int callCount = entry.getValue();
294             String ancestor = leafPkg;
295             while (ancestor != null && allPkgNames.contains(ancestor)) {
296                 if (ancestor.equals(toPkg)) break;
297                 weights.get(ancestor).merge(toPkg, callCount, Integer::sum);
298                 ancestor = extractPackageName(ancestor);
299             }
300         }
301     }
302     return weights;
303 }
304
305 // -----
306 // Step 5 -- reverse dependencies
307 // -----

```

```

308
309     private void updateDependentRelationships(DomainModel model) {
310         for (DomainModel.CalculatedElementInfo cls : model.getAllClasses().values()) {
311             for (String dep : cls.dependencies) {
312                 DomainModel.CalculatedElementInfo d = model.getClass(dep);
313                 if (d != null) d.addDependent(cls.fullName);
314             }
315         }
316         for (DomainModel.CalculatedElementInfo pkg : model.getAllPackages().values()) {
317             for (String dep : pkg.dependencies) {
318                 DomainModel.CalculatedElementInfo d = model.getPackage(dep);
319                 if (d != null) d.addDependent(pkg.fullName);
320             }
321         }
322     }
323
324     // -----
325     // Utilities
326     // -----
327
328     private static boolean isInSameSubtree(String a, String b) {
329         return a.startsWith(b + ".") || b.startsWith(a + ".");
330     }
331
332     private static String extractPackageName(String className) {
333         if (className == null || !className.contains(".")) return null;
334         return className.substring(0, className.lastIndexOf('.'));
335     }
336 }

```

The companion class `LocalLayerCalculator` implements Phase 3 (Step 6).

```

1 package de.weigend.s202.domain.architecture;
2
3 import de.weigend.s202.domain.DomainModel;
4 import de.weigend.s202.domain.DomainModel.CalculatedElementInfo;
5 import de.weigend.s202.graph.StronglyConnectedComponent;
6 import de.weigend.s202.graph.TarjanSCCFinder;
7 import de.weigend.s202.reader.DependencyModel;
8
9 import java.util.ArrayList;
10 import java.util.HashMap;
11 import java.util.HashSet;
12 import java.util.LinkedHashSet;
13 import java.util.List;
14 import java.util.Map;
15 import java.util.Set;
16
17 /**
18  * Assigns a per-parent layer position ({@code localLayerIndex}) to every
19  * element in a {@link DomainModel}. Each parent package is treated as a
20  * self-contained universe: only class dependencies whose source <em>and</em>
21  * target both live in the parent's subtree contribute to the sibling-only
22  * weighted graph. Refs that leave the parent are ignored -- they belong to
23  * the global {@code architectureLevel} via the package hierarchy.
24  *
25  * <p>Within each parent's sibling graph:
26  * <ol>
27  *   <li>Tarjan finds SCCs.</li>
28  *   <li>SCCs of size > 1 are broken by weight-based rank: a back-edge
29  *       leaves the smaller-rank member, the dominant direction stays.</li>
30  *   <li>Longest-path on the SCC-collapsed DAG assigns a {@code layer}

```

```

31 *      (= local layer index) per sibling. Layer 0 = bottom of the box,
32 *      higher values = visually higher.</li>
33 * </ol>
34 *
35 * <p>Single-child containers stay at layer 0 -- no graph, no work.
36 *
37 * <p>The algorithm intentionally mirrors the structure of
38 * {@link LevelCalculator}'s package-level computation, but applied
39 * locally and using a different graph. Local cycles are not surfaced
40 * as tangles; the user-visible tangle list remains the global one.
41 */
42 public class LocalLayerCalculator {
43
44     private static final double RANK_THRESHOLD = 0.1;
45
46     public void assign(DomainModel domain, DependencyModel rawModel) {
47         Map<String, List<CalculatedElementInfo>> childrenByParent = groupChildrenByParent(domain);
48         for (Map.Entry<String, List<CalculatedElementInfo>> entry : childrenByParent.entrySet()) {
49             assignForParent(entry.getValue(), domain, rawModel);
50         }
51     }
52
53     private void assignForParent(List<CalculatedElementInfo> siblings,
54                                 DomainModel domain, DependencyModel rawModel) {
55         if (siblings.size() <= 1) {
56             return; // single child stays at layer 0
57         }
58
59         Set<String> siblingFqns = new LinkedHashSet<>();
60         for (CalculatedElementInfo c : siblings) {
61             siblingFqns.add(c.fullName);
62         }
63
64         Map<String, Map<String, Integer>> weights = buildSiblingGraph(siblingFqns, domain, rawModel);
65         Map<String, Integer> layers = computeLayers(weights, siblingFqns);
66         for (CalculatedElementInfo s : siblings) {
67             s.setLocalLayerIndex(layers.getOrDefault(s.fullName, 0));
68         }
69     }
70
71     /**
72      * Build the weighted sibling-only edge map. For every class C in any
73      * sibling's subtree, every dep into another sibling's subtree
74      * contributes {@code callCount(C, dep)} weight; deps that leave the
75      * parent's subtree or stay inside the same sibling are skipped.
76      */
77     private Map<String, Map<String, Integer>> buildSiblingGraph(Set<String> siblingFqns,
78                                                                 DomainModel domain,
79                                                                 DependencyModel rawModel) {
80         Map<String, Map<String, Integer>> weights = new HashMap<>();
81         for (String s : siblingFqns) {
82             weights.put(s, new HashMap<>());
83         }
84         for (CalculatedElementInfo cls : domain.getAllClasses().values()) {
85             String fromSibling = containingSibling(cls.fullName, siblingFqns);
86             if (fromSibling == null) {
87                 continue;
88             }
89             for (String dep : cls.dependencies) {
90                 if (domain.getClass(dep) == null) {
91                     continue; // external -- not in any sibling
92                 }

```

```

93         String toSibling = containingSibling(dep, siblingFqns);
94         if (toSibling == null || toSibling.equals(fromSibling)) {
95             continue; // out of parent OR intra-sibling
96         }
97         int w = callCount(cls.fullName, dep, rawModel);
98         weights.get(fromSibling).merge(toSibling, w, Integer::sum);
99     }
100 }
101 return weights;
102 }
103
104 /**
105  * Walk up the fqcn package chain until the first ancestor that is one
106  * of the siblings, or return {@code null} when none of the ancestors
107  * is -- the class lives outside the parent's subtree.
108  */
109 private static String containingSibling(String fqcn, Set<String> siblingFqns) {
110     String current = fqcn;
111     while (current != null && !current.isEmpty()) {
112         if (siblingFqns.contains(current)) {
113             return current;
114         }
115         int dot = current.lastIndexOf('.');
116         if (dot < 0) {
117             return null;
118         }
119         current = current.substring(0, dot);
120     }
121     return null;
122 }
123
124 private static Map<String, List<CalculatedElementInfo>> groupChildrenByParent(DomainModel domain) {
125     Map<String, List<CalculatedElementInfo>> result = new HashMap<>();
126     for (CalculatedElementInfo cls : domain.getAllClasses().values()) {
127         result.computeIfAbsent(parentOf(cls.fullName), k -> new ArrayList<>()).add(cls);
128     }
129     for (CalculatedElementInfo pkg : domain.getAllPackages().values()) {
130         result.computeIfAbsent(parentOf(pkg.fullName), k -> new ArrayList<>()).add(pkg);
131     }
132     return result;
133 }
134
135 private static String parentOf(String fqcn) {
136     int dot = fqcn.lastIndexOf('.');
137     return dot < 0 ? "" : fqcn.substring(0, dot);
138 }
139
140 /**
141  * Method-call count from class {@code from} to class {@code to}.
142  * Falls back to 1 for structural deps (extends/implements/field type)
143  * where no method-call info exists -- same convention the global
144  * weighted package graph uses.
145  */
146 private static int callCount(String from, String to, DependencyModel rawModel) {
147     DependencyModel.ClassInfo cls = rawModel.getClass(from);
148     if (cls == null) {
149         return 1;
150     }
151     int count = 0;
152     String prefix = to + ".";
153     for (DependencyModel.MethodInfo method : cls.methods.values()) {
154         for (Map.Entry<String, Integer> call : method.methodCalls.entrySet()) {

```

```

155         if (call.getKey().startsWith(prefix)) {
156             count += call.getValue();
157         }
158     }
159 }
160 return count > 0 ? count : 1;
161 }
162
163 // ----- SCC-collapsed longest path with rank-based break -----
164
165 private Map<String, Integer> computeLayers(Map<String, Map<String, Integer>> weights,
166                                         Set<String> nodes) {
167     Map<String, Set<String>> graph = new HashMap<>();
168     for (String n : nodes) {
169         graph.put(n, new HashSet<>(weights.getOrDefault(n, Map.of()).keySet()));
170     }
171
172     // Identify and remove back-edges within SCCs based on weight rank.
173     boolean changed = true;
174     while (changed) {
175         changed = false;
176         for (StronglyConnectedComponent scc : new TarjanSCCFinder(graph).findSCCs()) {
177             if (scc.getSize() < 2) {
178                 continue;
179             }
180             Set<String> members = scc.getMembers();
181             Map<String, Double> rank = new HashMap<>();
182             for (String m : members) {
183                 int out = 0;
184                 int in = 0;
185                 for (String other : members) {
186                     if (other.equals(m)) {
187                         continue;
188                     }
189                     out += weights.getOrDefault(m, Map.of()).getOrDefault(other, 0);
190                     in += weights.getOrDefault(other, Map.of()).getOrDefault(m, 0);
191                 }
192                 rank.put(m, (out - in) / (double) Math.max(1, out + in));
193             }
194             for (String from : new ArrayList<>(members)) {
195                 for (String to : new ArrayList<>(graph.getOrDefault(from, Set.of()))) {
196                     if (!members.contains(to)) {
197                         continue;
198                     }
199                     if (rank.get(from) < rank.get(to) - RANK_THRESHOLD) {
200                         graph.get(from).remove(to);
201                         changed = true;
202                     }
203                 }
204             }
205         }
206     }
207
208     // SCC-collapsed longest-path level assignment.
209     List<StronglyConnectedComponent> sccs = new TarjanSCCFinder(graph).findSCCs();
210     Map<String, StronglyConnectedComponent> nodeToScc = new HashMap<>();
211     for (StronglyConnectedComponent scc : sccs) {
212         for (String m : scc.getMembers()) {
213             nodeToScc.put(m, scc);
214         }
215     }
216     Map<Integer, Set<Integer>> sccDeps = new HashMap<>();

```

```

217     for (StronglyConnectedComponent scc : sccs) {
218         sccDeps.put(scc.getId(), new HashSet<>());
219         for (String m : scc.getMembers()) {
220             for (String to : graph.getDefault(m, Set.of())) {
221                 StronglyConnectedComponent toScc = nodeToScc.get(to);
222                 if (toScc != null && toScc.getId() != scc.getId()) {
223                     sccDeps.get(scc.getId()).add(toScc.getId());
224                 }
225             }
226         }
227     }
228     Map<Integer, Integer> sccLayers = new HashMap<>();
229     for (StronglyConnectedComponent scc : sccs) {
230         sccLayers.put(scc.getId(), 0);
231     }
232     boolean lvlChanged = true;
233     while (lvlChanged) {
234         lvlChanged = false;
235         for (StronglyConnectedComponent scc : sccs) {
236             int maxDep = -1;
237             for (int depId : sccDeps.get(scc.getId())) {
238                 maxDep = Math.max(maxDep, sccLayers.get(depId));
239             }
240             int newLayer = maxDep >= 0 ? maxDep + 1 : 0;
241             if (sccLayers.get(scc.getId()) != newLayer) {
242                 sccLayers.put(scc.getId(), newLayer);
243                 lvlChanged = true;
244             }
245         }
246     }
247
248     Map<String, Integer> result = new HashMap<>();
249     for (StronglyConnectedComponent scc : sccs) {
250         int layer = sccLayers.get(scc.getId());
251         for (String m : scc.getMembers()) {
252             result.put(m, layer);
253         }
254     }
255     return result;
256 }
257 }

```

Open Points Before Submission

- Keep the submission abstract to one page in the INSTICC template.